

# **Data Mining**

## **- Deep Learning, Transformers, and Transfer Learning**

---

Joe Burdis  
Fall 2024  
CUNY Graduate Center

# Outline

1. HW2 discussion
2. Week 6 recap
3. Deep Learning
  - a. CNNs (for computer vision)
  - b. Audio
  - c. Transformers for NLP
4. using BERT (or BERT-like models)
  - a. <https://huggingface.co/docs/transformers/notebooks>
  - b. [https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify\\_text\\_with\\_bert.ipynb](https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify_text_with_bert.ipynb)
5. General discussion of transfer learning and fine-tuning

# Docker Tutorial

For next week:

<https://docs.docker.com/language/python/build-images/>


# BM25 vs. Embedding Models

- BM25: Bag-of-words model, term frequency focus.
- Embeddings: Dense vector representations capturing semantic meaning.
- BM25: Efficient for large-scale search, interpretable.
- Embeddings: Better at capturing context and synonyms.
- BM25: Requires careful tuning of parameters (defaults aren't bad though)
- Embeddings: Pre-trained models available, reducing training needs.
- Hybrid approaches can leverage strengths of both.

# The Biological Neuron


The behavior occurs at the membrane between the inside and outside of the cell.

- Dendrites: appendages that are designed to receive communications via **impulses** from other cells and transmit them to the cell body.
- Soma (or cell body): combines the impulses received and triggers action on the axons.
- Axons: the long thread-like part of a neuron that impulses are sent to other cells.



When the synapse releases neurotransmitters, it opens some sodium channels. If **ENOUGH** channels are opened, a region of the membrane will depolarize. **Action potential! So whether an AP occurs or not is a sum!**

Neuron Networks:



# (Artificial) Neural Networks

The initial neural network is often cited as Rosenblatt's Perceptron for binary classification which consisted of input values, weights (and bias), a weighted sum, and activation function. "Perceptron" and "neuron" are often used interchangeably, but "perceptron" refers to something more specific (so use "neuron").

Some basic features of Neural Networks

1. Each activation is differentiable (usually nonlinear)
2. Contains hidden layers (non-input or output layer)
3. High degree of connectivity (E.g. "Fully Dense")

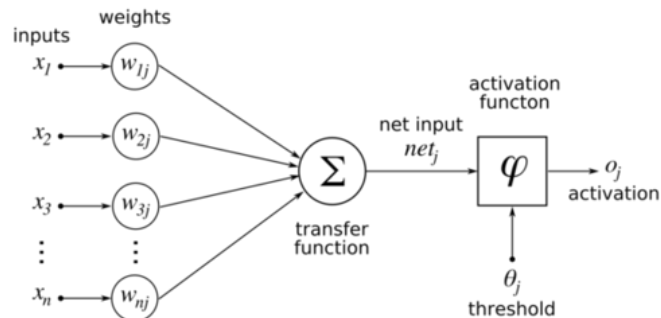
Note: "deep" refers to more than 1 hidden layer

Activation function for neuron:


$$h_{ij}(x) = \phi(w_{ij}^T x + b_{ij})$$

$\phi: \mathbb{R} \rightarrow \mathbb{R}$

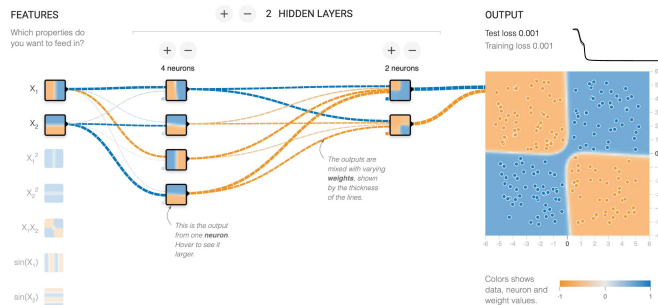
Rosenblatt's Perceptron:



Neural Network:



# Universal Approximation Theorem



*“A neural network with one hidden layer containing a sufficient but finite number of neurons can approximate any continuous function to a reasonable accuracy, under certain conditions for activation functions”*

*Let  $\varphi(\cdot)$  be a nonconstant, bounded, and monotone-increasing continuous function. Let  $I_{m_0}$  denote the  $m_0$ -dimensional unit hypercube  $[0, 1]^{m_0}$ . The space of continuous functions on  $I_{m_0}$  is denoted by  $C(I_{m_0})$ . Then, given any function  $f \in C(I_{m_0})$  and  $\varepsilon > 0$ , there exist an integer  $m_1$  and sets of real constants  $\alpha_i$ ,  $b_j$ , and  $w_{ij}$ , where  $i = 1, \dots, m_1$  and  $j = 1, \dots, m_0$  such that we may define*

$$F(x_1, \dots, x_{m_0}) = \sum_{i=1}^{m_1} \alpha_i \varphi \left( \sum_{j=1}^{m_0} w_{ij} x_j + b_i \right) \quad (4.88)$$

*as an approximate realization of the function  $f(\cdot)$ ; that is,*

$$|F(x_1, \dots, x_{m_0}) - f(x_1, \dots, x_{m_0})| < \varepsilon$$


*for all  $x_1, x_2, \dots, x_{m_0}$  that lie in the input space.*

*So just use neural networks for everything, right?*

**No!**

Performance

average



**There is No Free Lunch**

# What are Neural Networks used for?

## Computer Vision (CV) / Image Processing

- Image classification
- Facial recognition <Optical Character Recognition(OCR)>
- Object detection
- Image (and semantic) segmentation
- Any model using images as inputs (text featurizers)

## Tabular Data

- Model using an arrangement of information or data, typically in rows and columns
- Historically only good for very large data and rare instances, but recent innovations have proved quite useful

## Natural Language Processing (NLP)

- Translation
- Sentiment analysis
- Text summarization or extraction (NER)
- Topic modeling
- Any model using text as inputs (text featurizers)

## Other

- Reinforcement Learning (DQN)
- Generative models
- Knowledge graphs
- Art, Music, Gaming
- Dimensionality reduction / compression



# Why Neural Networks and why “deep”?

## Pros:

- NNs can outperform other ML techniques on large and complex problems and are SOTA in many modalities (images, audio, text, etc.)
- The tremendous increase in **data availability and** computing power has made it possible to train large neural networks in a reasonable amount of time.
- The training algorithms have been improved
  - Activation functions: Relu
  - Heuristics: batch normalization and learning rate scheduling/annealing
  - Architectures: CNN, GRU/LSTM, Transformers/Attention, and others
- Some theoretical limitations of NNs have turned out to be benign in practice.

## Cons:

- Huge number of parameters: slow to use, prone to overfitting.
- Black box model: hard to understand the learned function
- High Variance and lack of performance guarantee

## Mitigations:

- **Proper experimentation techniques!!!**
- **Use good baselines!!!**
- Early stopping (careful / double descent)
- New explanation methods
  - Activations
  - Integrated gradients
  - Model agnostic methods / testing

**Let's go through XOR on the board!**

# Backpropagation

1. **Initialization of weights (and biases):** Choose weights randomly(see initializers)
  - a. Range of the weights should be determined so the induced fields of the first hidden layer are near the midrange of the activation function.
  - b. **Note:** it is common to adjust the range of the input vectors have magnitudes of the input terms in the range  $[-1, 1]$  or standardized (aligns with adjustment of the initial weights and any scaling parameters)
2. **Forward computation / propagation (or pass):** Present  $\mathbf{x}_i$  to the input and  $d_i$  at the output to compute error.
3. **Backward computation / propagation:** Compute local gradients.

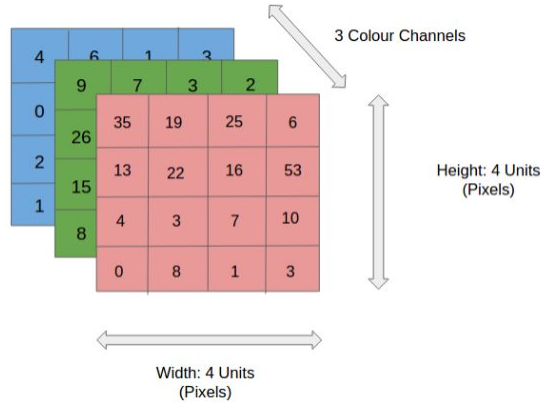
The math:

<https://www.youtube.com/watch?v=tleHLnjs5U8>

The intuition:

<https://www.youtube.com/watch?v=llg3gGewQ5U>

# Computer Vision




CNN Explainer: <https://poloclub.github.io/cnn-explainer/>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

[https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning)

# Audio Data



<https://towardsdatascience.com/audio-deep-learning-made-simple-sound-classification-step-by-step-cebc936bbe5>

# Multi-Head Attention

Computing  $e_1, \dots, e_n \in \mathbb{R}$  from  $k_1, \dots, k_n \in \mathbb{R}^{d_1}$  and  $q \in \mathbb{R}^{d_2}$

- Basic dot-product attention:

$$e_i = q^T k_i \in \mathbb{R}$$

where  $d_1 = d_2$

- Multiplicative attention:

$$e_i = q^T W k_i \in \mathbb{R}$$

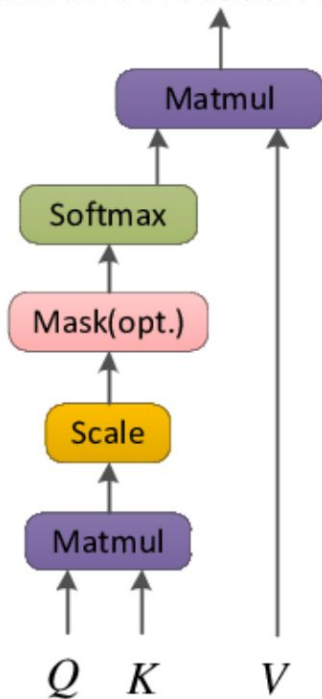
where  $W \in \mathbb{R}^{d_2 \times d_1}$  is a weight matrix

- Additive attention:


$$e_i = W_3^T \tanh(W_1 k_i + W_2 q) \in \mathbb{R}$$

where  $W_1 \in \mathbb{R}^{d_3 \times d_1}$ ,  $W_2 \in \mathbb{R}^{d_3 \times d_2}$  are weight matrices and  $W_3 \in \mathbb{R}^{d_3}$  is a weight vector.

Scaled Dot-Product Attention



Multi-Head Attention

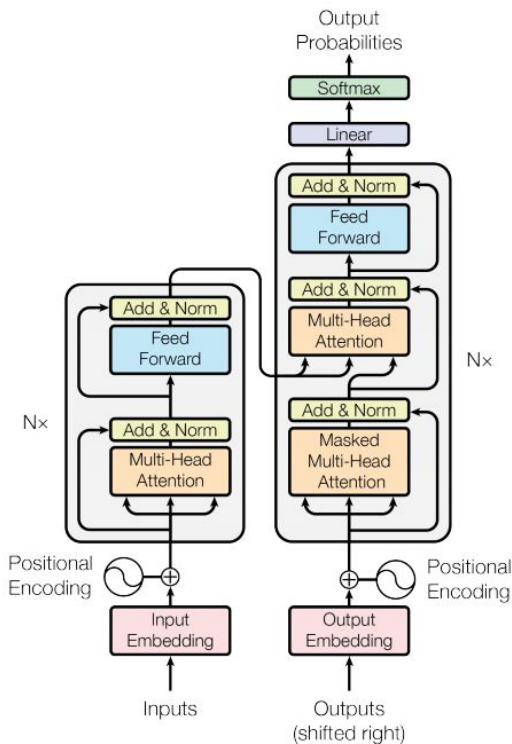


# Transformer (used in NLP and other tasks)

Multiple heads and multiple transformer layer / blocks

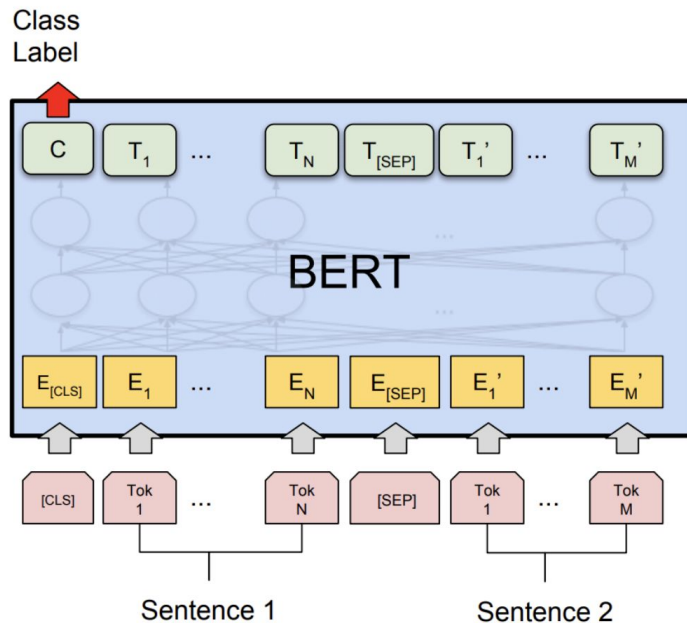
**Left:** Encoder (BERT)

**Right:** Decoder (e.g. GPT)



# NLP with BERT

Encode Only



[https://pytorch.org/tutorials/intermediate/dynamic\\_quantization\\_bert\\_tutorial.html](https://pytorch.org/tutorials/intermediate/dynamic_quantization_bert_tutorial.html)

[https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify\\_text\\_with\\_bert.ipynb](https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify_text_with_bert.ipynb)

# GPT Architecture

Decoder Only

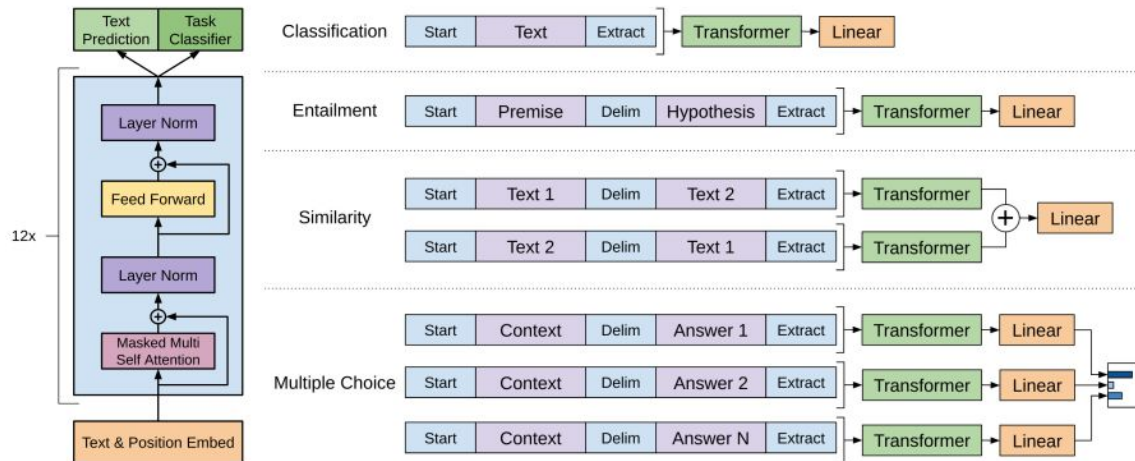
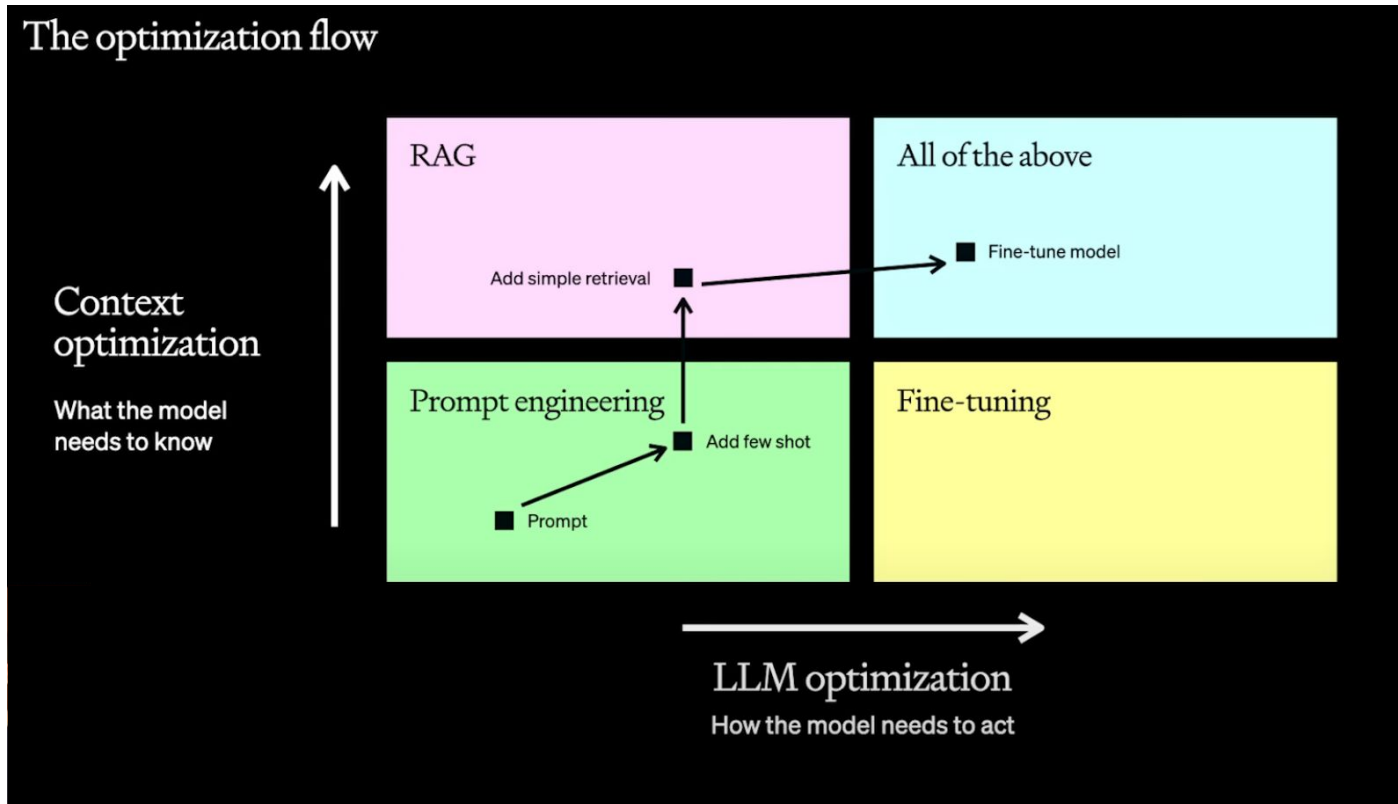


Figure 1: **(left)** Transformer architecture and training objectives used in this work. **(right)** Input transformations for fine-tuning on different tasks. We convert all structured inputs into token sequences to be processed by our pre-trained model, followed by a linear+softmax layer.

<https://paperswithcode.com/paper/improving-language-understanding-by>



# Improving LLM Pipelines



# TikTok Recommendation System

<https://analyticsindiamag.com/tiktok-parent-bytedance-reveals-its-sota-recommendation-engine/>

<https://arxiv.org/pdf/2209.07663.pdf>

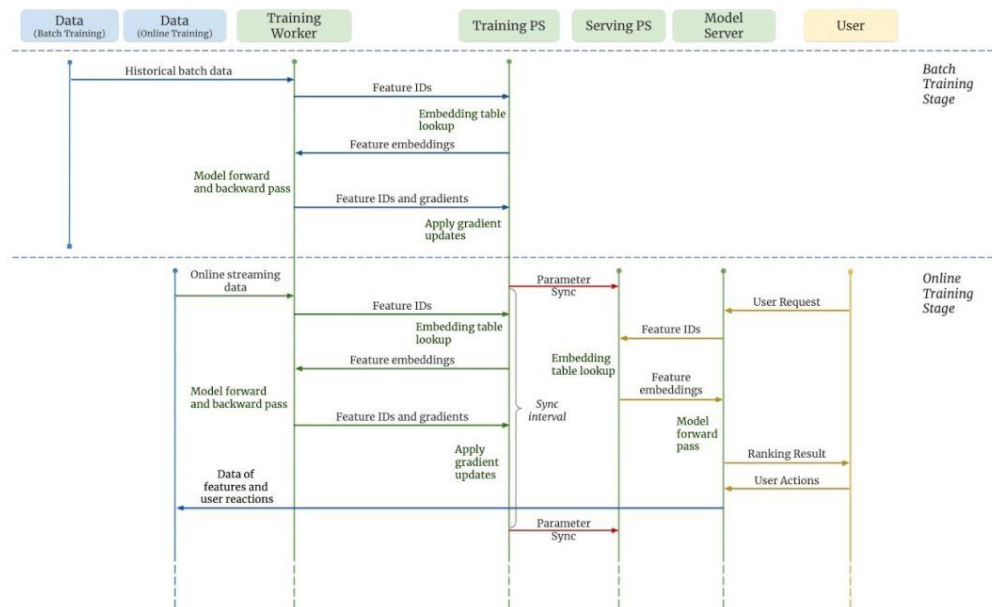
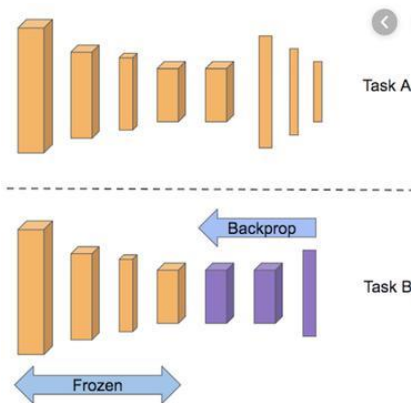


Figure 1: Monolith Online Training Architecture.

# Transfer Learning

- One can think of deep learning models as learning a complex feature extractor and using that to predict a given target.
- **Issue:** In order to learn this feature extractor, we need lots of data to train from scratch.
- What about using the first X layers of a model trained on a task with lots of data and using those trained weights for another task?



( Image credit: [Subodh Malgonde](#) )

# Transfer Learning

- You need less data - only need to train purple layers
- Typically works very well if new task in same domain as original task
- **Example:** ImageNet or other networks trained on ImageNet dataset may fail if using strict transfer learning to detect human emotions (no faces in dataset)
- **Solution:** Fine-tuning - i.e. the layers borrowed don't need to be frozen

**CNN Example:** [https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning)

**NLP Example:**

[https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify\\_text\\_with\\_bert.ipynb](https://colab.research.google.com/github/tensorflow/text/blob/master/docs/tutorials/classify_text_with_bert.ipynb)

**Hugging Face:** <https://huggingface.co/docs/transformers/notebooks>