

Data Mining Recommendation Systems



Joseph Burdis
Fall 2024
CUNY Graduate Center

Topic List

- Formulation of the recommendation system
- Content-based approach
- Collaborative filter approach

Problem Formulation

Recommendation System

- Amazon
- Spotify
- Netflix
- Youtube
- Google News
- ...

Recommendation Matters

In 1988, a British mountain climber named Joe Simpson wrote a book called *Touching the Void*, a harrowing account of near death in the Peruvian Andes. It only achieved a modest success and was soon forgotten.

A decade later, an American writer and mountaineer Jon Krakauer wrote a best-selling book *Into Thin Air*, another book about a mountain-climbing tragedy. It caused *Touching the Void* to sell again. The book was on the New York Time bestseller list for 14 weeks.

It was Amazon recommendation who suggested *Touching the Void* to readers of *Into Thin Air*. The latter book was actually out-of-print then. A new edition has to be published to meet readers' request.

Types of Recommendations

Editorial and hand curated

- List of favorites
- List of “essential” items

Simple aggregates

- Top 10
- Most popular

Tailored to individual users

- Amazon, Netflix, ...

Formal Model: Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Key Problems

- How to collect data in the utility matrix?
 - Explicit: ask people to rate items
 - Implicit: learn ratings from user actions
- How to extrapolate unknown ratings from the known ones?
 - Most people have not rated most items
 - Cold start: new item, new user
- How to measure the performance of extrapolation methods?

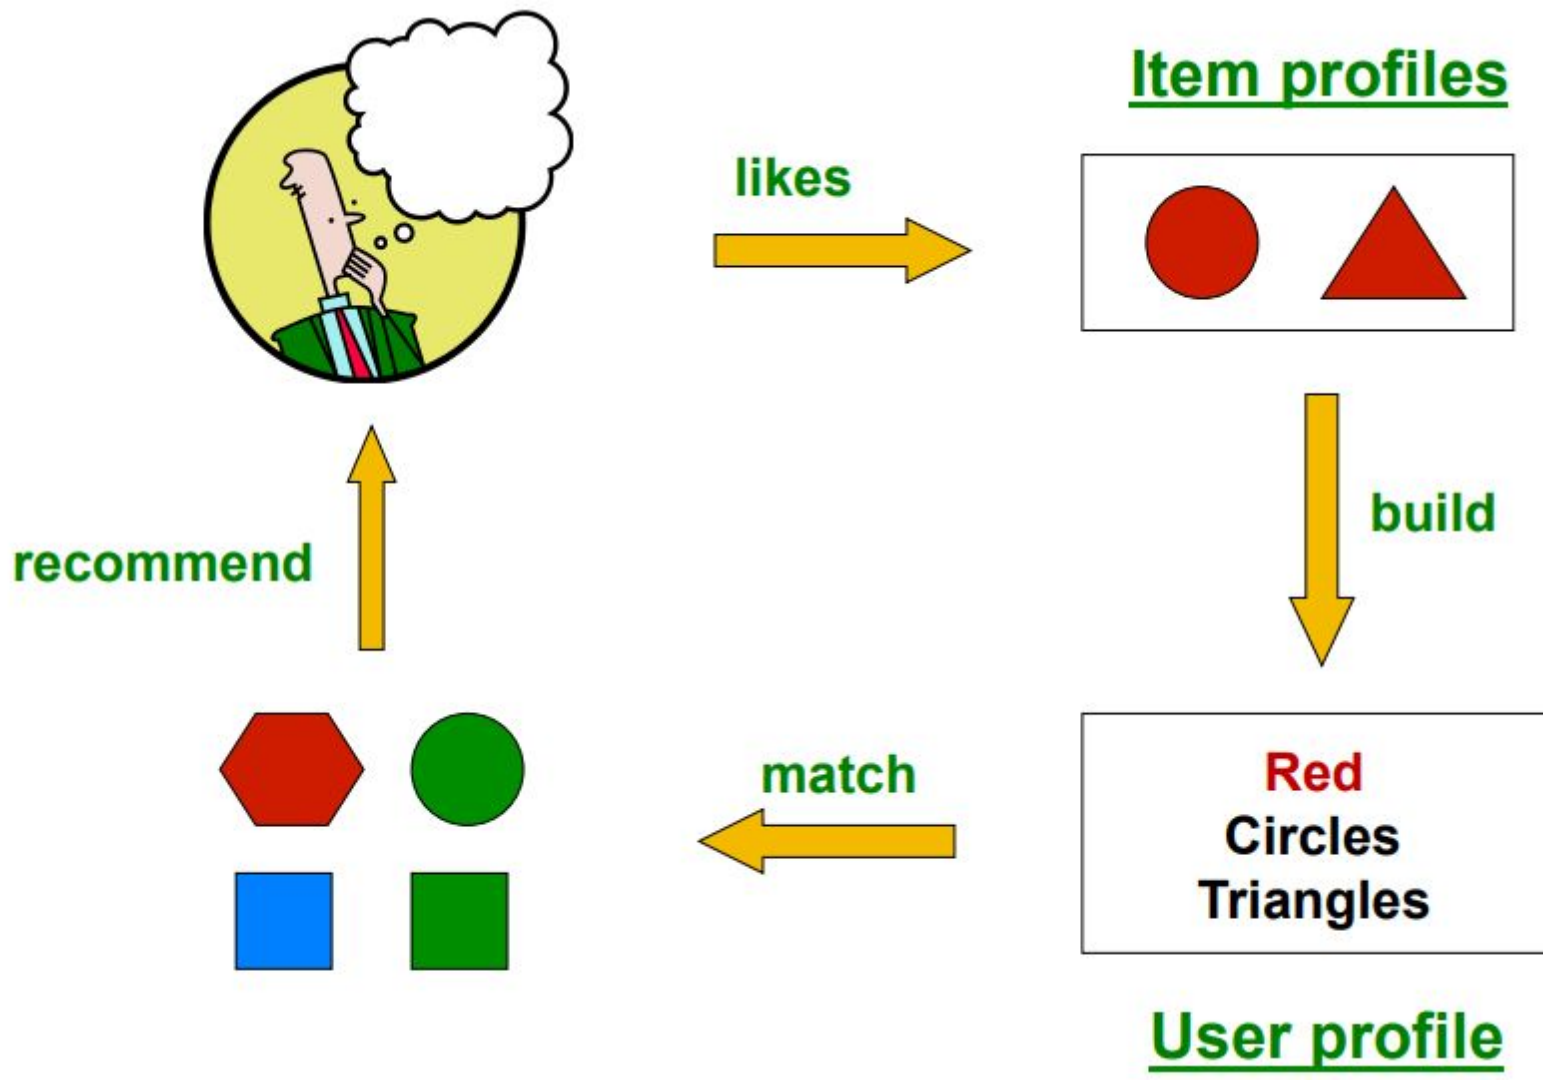
Content-Based Recommendations

Content-Based Recommendations

Main Idea: Recommend items to user x similar to previous items rated highly by x .

Challenge: How to measure similarity?

- Movie recommendations: same genre, same actors, same director, ...
- Websites: similar topic, similar content, similar opinion, ...



Item Profiles

Content-based approach require an **item profile** for each item.

Item profile is a vector of features

- Movies: author, title, actors, director, ...
- Text: topic, keywords, ...

Challenge:

- How to effectively represent an item?
- How to identify important features?

Example: TF-IDF

For text mining, it is common to use TF-IDF to measure the importance of words.

TF-IDF: Term Frequency * Inverse Document Frequency

- If a word appears frequently in a document, it is likely to be important.
- If a word rarely appears in a random document but it appears here, it is likely to be important.

More precisely:

- TF: frequency of word / maximum word frequency in document
- IDF: $\log(\text{total number of documents} / \text{number of docs that has this word})$
- TF-IDF score: $\text{TF} * \text{IDF}$
- Usage: pick words with highest TF-IDF scores

User Profile

User profile can be created by averaging the related item profiles

- Weighted average of rated item profiles
- Or: weight by difference from average rating for each item

Measure similarity:

- Cosine distance: Given user profile x and item profile i :
- $u(x, i) = \cos(x, i) = x \cdot i / (\|x\| \cdot \|i\|)$

Pros and Cons

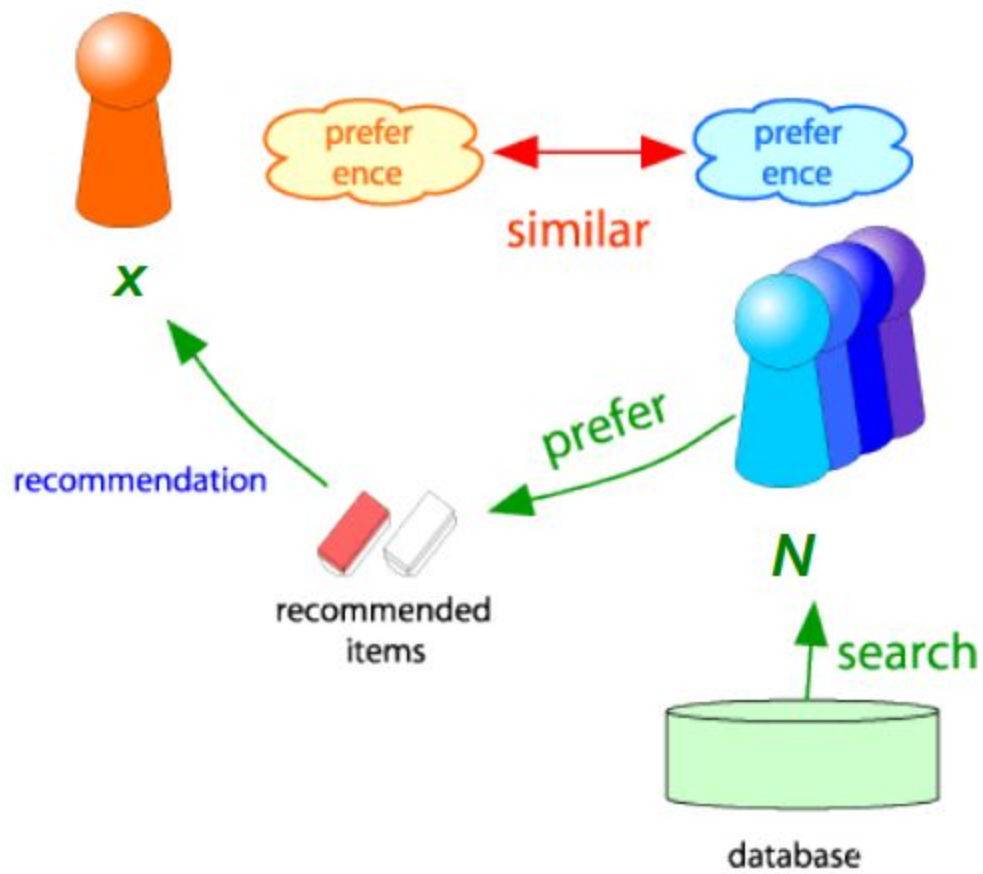
- + No need for data on other users
 - + Able to recommend to users with unique tastes
 - + Able to recommend new & unpopular items
 - + Able to provide explanations
-
- Finding the appropriate features is challenging
 - Hard to build profile for new users
 - Never recommend items outside user's content profile
 - Averaging of multiple user interests may result error
 - Unable to exploit quality judgments of other users

Collaborative Filtering

Collaborative Filtering

Main idea: Find recommendations by looking at similar users

- Consider user x
- Find set N of other users whose ratings are similar to x 's ratings
- Estimate x 's ratings based on ratings of these N users
- Recommend items with potential high rating



Challenge: Finding Similar Users

- Represent each user's ratings as a vector
- Jaccard similarity = $|\text{intersection}| / |\text{union}|$
 - Ignores value of the ratings
 - Only considers overlap of rated items
- Cosine similarity
 - Missing ratings have similar impacts as negative ratings
- Pearson correlation coefficient
- Matrix decomposition

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3					3

- **Intuitively we want:** $\text{sim}(A, B) > \text{sim}(A, C)$
- **Jaccard similarity:** $1/5 < 2/4$
- **Cosine similarity:** $0.386 > 0.322$
 - Considers missing ratings as “negative”

■ **Solution: subtract the (row) mean**

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	$2/3$			$5/3$	$-7/3$		
B	$1/3$	$1/3$	$-2/3$				
C				$-5/3$	$1/3$	$4/3$	
D		0					0

sim A,B vs. A,C:
 $0.092 > -0.559$

Notice cosine sim. is correlation when data is centered at 0

Rating Predictions

Let r_x be the vector of user x 's ratings. Let N be the set of k users most similar to x .

Predicting user x 's rating for item i :

- Option 1:

$$r_{xi} = 1/k \sum_{y \in N} r_{yi}$$

- Option 2:

$$r_{xi} = \sum_{y \in N} s_{xy} r_{yi} / \sum_{y \in N} s_{xy}$$

Item-Item Collaborative Filtering

So far we have been discussing user-user collaborative filtering. In practice, item-item filtering is more successful

- Each item in general receives a large amount of ratings
- Items are simpler than users to measure
- Item similarity is more meaningful than user similarity
- The nature of an item stays constant over time

The computation model is similar to user-user model.

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3		?	5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	



- estimate rating of movie 1 by user 5

Example: Estimate Movie Rating

- Suppose $N = 2$
- Use Pearson correlation as similarity
- Use weighted average to predict movie rating

	Users											
	1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1	3		2.6	5			5		4	
	2		5	4			4			2	1	3
	<u>3</u>	2	4	1	2		3		4	3	5	
	4	2	4		5			4			2	
	5		4	3	4	2					2	5
	<u>6</u>	1	3		3			2			4	

Predict by taking weighted average:

$$r_{1.5} = (0.41*2 + 0.59*3) / (0.41+0.59) = 2.6$$

Matrix Factorization

Matrix factorization is a simple embedding model. Given the feedback matrix $A \in \mathbb{R}^{m \times n}$, where m is the number of users (or queries) and n is the number of items, the model learns:

- A user embedding matrix $U \in \mathbb{R}^{m \times d}$, where row i is the embedding for user i .
- An item embedding matrix $V \in \mathbb{R}^{n \times d}$, where row j is the embedding for item j .



Matrix Factorization

Observed Only MF

1		1	1	
	1			1
1	1	1		
			1	1

$$\sum_{(i,j) \in \text{obs}} (A_{ij} - U_i \cdot V_j)^2$$

Weighted MF

1	0	1	1	0
0	1	0	0	1
1	1	1	0	0
0	0	0	1	1

$$\sum_{(i,j) \in \text{obs}} (A_{ij} - U_i \cdot V_j)^2 + w_0 \sum_{(i,j) \notin \text{obs}} (0 - U_i \cdot V_j)^2$$

SVD

1	0	1	1	0
0	1	0	0	1
1	1	1	0	0
0	0	0	1	1

$$\|A - UV^T\|_F^2 = \sum_{(i,j)} (A_{ij} - U_i \cdot V_j)^2$$

Pros and Cons

- + No feature selection needed
- Cold start: need enough users in the system to find a match
- Sparsity: rating matrix is sparse
- First rater: new items has no ratings
- Popularity bias: tend to recommend popular items
- Cannot recommend items to someone with unique taste